

# EUfficient Reachability in Software with Arrays

Denis Bueno

Computer Science and Engineering  
University of Michigan  
Email: dlbueno@umich.edu

Arlen Cox

Institute for Defense Analyses  
Center for Computing Sciences  
Email: arlen@super.org

Karem Sakallah

Computer Science and Engineering  
University of Michigan  
Email: karem@umich.edu

**Abstract**—Whether representing strings, heap objects, or numerical vectors, arrays are pervasive in software. Unfortunately, while several software model checkers support arrays, they tend to struggle with many array-manipulating programs due to work expended generating theory lemmas that are ultimately irrelevant or redundant. By judicious abstraction of array operations to the logic of equality with uninterpreted functions (EUF), we show that we can *directly* reason about array reads and *adaptively* learn lemmas about array writes leading to significant performance improvements over existing approaches. We find that our model checker solves more than 100 more SV-COMP benchmarks than SPACER, a leading model checker.

## I. INTRODUCTION

Arrays and array-like structures are pervasive in the software world. From C/C++ arrays and vectors to Python lists, it is difficult to find software that doesn't use and manipulate arrays. Despite this, research of software model checkers has largely focused on finding numerical invariants and proving numerical properties of programs. As results of the software verification competition (SV-COMP) show, even when model checkers support arrays, there are a significant number of programs that cannot be automatically verified—some for a lack of expressivity and some for a lack of performance. Our focus is on the latter.

The key challenge that we face is adequately controlling theory reasoning in the SMT solver underlying the model checker. While SMT solvers typically have an array theory and can therefore directly solve array problems, the interface that SMT solvers provide does not provide for adequate incrementality and hinting to enable maximal performance. For instance, we find that, in SV-COMP benchmarks, as many as 90% of the array lemmas that the SMT solver is learning are either redundant or ultimately irrelevant. Most lemmas either do not advance the cause of the model checker or were thrown away by the SMT solver due to imperfect caching. Thus time spent learning those lemmas was wasted effort.

To eliminate this waste, we do incremental inductive model checking on top of an equality with uninterpreted functions (EUF) theory [1]. This removes need for SMT array theories in the core incremental model checking process, relegating the array theory solely to abstraction refinement operations. By eliminating arrays from the step-wise refinement operations, the number of operations that can do redundant or irrelevant work is reduced by a factor of thousands. Additionally this means that array lemmas are only learned where they are pertinent to proving or disproving the property.

Moreover our strategy addresses a fundamental tension. On the one hand, incremental model checkers [2], which construct a safety proof bit by bit, are particularly scalable because their many individual queries are simple to solve and generalize. On the other hand, these queries lack error path information that could simplify overall checking.

For example, consider model checking the following program, assuming that  $a$ ,  $b$ , and  $f$  are distinct constant values:

```
int[] A; int i, a, b, f;
 $\ell_1$ : A[3] = f;
 $\ell_2$ : A[1] = a;
      A[2] = b;
      assume(1 <= i <= 3);
      if (A[i] == f);
 $\ell_3$ :   error();
      else
 $\ell_4$ :   exit();
```

The model checker is trying to find if any values of  $i$  lead to the error at location  $\ell_3$ . Of course it can reach  $\ell_3$  if  $i = 3$ , which the checker takes two SMT queries to discover. The first query corresponds to reaching  $\ell_3$ , where  $A[i] = f$ , from  $\ell_2$ . The solver deduces  $i \notin \{1, 2\}$ , meaning the property may yet be violated, so the checker moves on to the next query, which corresponds to reaching the failure from  $\ell_1$ . The first query involves two array stores and one read; the SMT array theory will generate theory lemmas to deduce that  $A[i]$  is not set to  $f$  by any assignment from  $\ell_2$ . Several of the lemmas ultimately do not matter, however, since the property is discovered to be violated by the antecedent assignment at  $\ell_1$ .

We study arrays and array abstraction in the context of EUF model checking and make the following contributions:

- 1) We develop an algorithm for integrating array abstraction into EUFORIA, an EUF-based, incremental, inductive, model checker (Section III & III-A).
- 2) We introduce a refinement procedure for learning relevant array lemmas (Section III-B & III-D).
- 3) We evaluate the integration of array abstraction with EUF-based model checking using a variety of device driver benchmarks from SV-COMP (Section IV). We find that EUFORIA performs well compared to SPACER, IC3IA, and ELDARICA.

Before introducing our approach, we discuss background.

## II. BACKGROUND

*a) Equality with Uninterpreted Functions (EUF):* We consider a first-order language with equality with signature  $\mathcal{S}$  and two common sorts, **BOOLS** and **INTS**. Our setting is standard quantifier-free, first-order logic (FOL) with the standard notions of theory, satisfiability, validity, entailment, and models.

The EUF logic grammar is presented here:

type		production	explanation
<i>term</i> ( $t$ )	$::=$	$x \mid y \mid z \mid \dots$	0-arity term
	$ $	$F(t_1, t_2, \dots, t_n)$	uninterp. function (UF)
	$ $	$\text{ite}(f, t_1, t_2)$	if-then-else
<i>atom</i> ( $a$ )	$::=$	$t_1 = t_2$	equality atom
	$ $	$x \mid y \mid z \mid \dots$	Boolean atom
	$ $	$P(t_1, t_2, \dots, t_n)$	uninterp. predicate (UP)
<i>formula</i> ( $f$ )	$::=$	$a$	
	$ $	$\neg a$	negation
	$ $	$f_1 \wedge f_2$	conjunction
	$ $	$f_1 \vee f_2$	disjunction

Atomic formulas (atoms) are made up of Boolean identifiers, uninterpreted predicates (UPs), and equalities between terms. Formulas are made up of terms combined with arbitrary Boolean structure. For simplicity, but without loss of generality, we only consider formulas in negation normal form. A *literal* is a (possibly-negated) atom containing no occurrences of ITE. A *clause* is a disjunction of literals. A *cube* is a conjunction of literals. When convenient, a formula  $F$  may be treated as a set of its top-level conjuncts, e.g.,  $x = 1 \in F$  if  $F = (x > 17 \wedge x = 1)$ .  $a \models b$  means that  $a$  entails  $b$ . We write uninterpreted objects—terms  $x$ , functions  $F$ , and predicates  $P$ —in sans serif face. The semantics of these formulas is standard.

*b) Arrays:* We consider a theory of arrays with extensionality and constant-initialized arrays. This theory has the particular function symbols `select`, `store`, and `const-array`. The theory is defined by McCarthy’s axioms [3], extended with axioms for extensionality and constant initialization:

$$\forall a i j e. i = j \implies \text{select}(\text{store}(a, i, e), j) = e \quad (1)$$

$$\forall a i j e. i \neq j \implies \text{select}(\text{store}(a, i, e), j) = \text{select}(a, j) \quad (2)$$

$$\forall a b i. a = b \implies \text{select}(a, i) = \text{select}(b, i) \quad (3)$$

$$\forall i k. \text{select}(\text{const-array}(k), i) = k \quad (4)$$

The first two axioms specify array accesses. The third axiom specifies that equal arrays have identical elements at identical indices. The fourth axiom specifies that every index of a constant-initialized array has the initializer value.

We consider this array theory—specifically including equality and constant initialization—because of its utility for software verification. Programs commonly bulk-initialize arrays and array equality allows encodings to be easily composed.

*c) Transition Systems for Programs:* A *transition system* [4], [5] is a tuple  $\mathcal{T} = (X, Y, I, T)$  consisting of a (non-empty) set of *state variables*  $X = \{x_1, \dots, x_n\}$ , a (possibly empty) set of *input variables*  $Y = \{y_1, \dots, y_m\}$ , and two formulas:  $I$ , the *initial states*, and  $T$ , the *transition relation*. Formulas over state variables, or *state formulas*, are identified with the sets of states they denote; for example, the formula

$(x_1 = x_2)$  denotes all states where  $x_1$  and  $x_2$  are equal, and other variables may have any value. The *state space* of  $\mathcal{T}$  is the set of all valuations to variables in  $X$ . The set of *next-state variables* is  $X' = \{x'_1, x'_2, \dots, x'_n\}$ . For a formula  $\sigma$ ,  $\text{Vars}(\sigma)$  denotes the set of state variables free in  $\sigma$  (respectively,  $\text{Vars}'(\sigma)$  denotes the set of next-state variables in  $\sigma$ ). We may write  $\sigma$  as  $\sigma(X)$  when we wish to emphasize that the free variables in  $\sigma$  are drawn solely from the set  $X$ , i.e.,  $\text{Vars}(\sigma(X)) \subseteq X$ ; similarly for  $\sigma(X')$  (also written  $\sigma'$ ). The system’s *transition relation*  $T(X, Y, X')$  is a formula over the current-state, next-state, and input variables.

We write a transition as  $\sigma(X) \xrightarrow{T} \omega(X)$  if each state in  $\sigma$  transitions to some state in  $\omega$  under  $T$ , i.e.,  $\sigma \wedge T \models \omega'$ . A (possibly-infinite) sequence of transitions  $\sigma_0(X) \xrightarrow{T} \sigma_1(X) \xrightarrow{T} \sigma_2(X), \dots$  is an *execution* of a transition system if  $\sigma_0(X) \models I(X)$ .

A *safety property* is specified by a predicate,  $P(X)$ . The *model checking problem* is to determine whether any state satisfying  $\neg P(X)$  is reachable through an execution of  $T$ . A *counterexample* to a safety property  $P(X)$  is a  $k$ -step execution such that  $\sigma_k(X) \models \neg P(X)$ .

A *concrete transition system* (CTS) is defined over bit vector and array state variables and operations in the quantifier-free logic of bit vectors and arrays (QF\_ABV from SMT-LIB [6]).

## III. MODEL CHECKING WITH EUF AND ARRAYS

This paper introduces a program abstraction using EUF that incorporates arrays. First, we review EUFORIA’s data abstraction approach, as it is the inspiration and basis for the array abstraction.

EUFORIA homomorphically maps bit vector operations into uninterpreted functions in order to avoid potentially expensive reasoning (e.g., nonlinear computations). EUF operation abstraction was introduced by Burch and Dill [7] for checking the equivalence between pipelined computer architectures and their single-step specifications. EUFORIA adopts and extends this abstraction to check for general safety properties. For purposes of this paper, we assume there is an abstraction function  $\llbracket \cdot \rrbracket$  that homomorphically maps a given concrete transition relation to an EUF transition relation. For instance,  $\llbracket x' = x + 1 \rrbracket = (\hat{x}' = \text{ADD}(\hat{x}, \hat{1}))$ . State variables, inputs, and constants are mapped to uninterpreted 0-arity terms with hats (e.g.,  $x \mapsto \hat{x}$ , and  $1 \mapsto \hat{1}$ ). Operations are mapped to appropriately-named UFs. The crucial property guaranteed by this abstraction is that executions of the EUF transition system over-approximate the executions of the concrete transition system. The details of EUFORIA’s translation are available in previous work [1].

This paper brings arrays into the mix. In order to avoid the overhead of instantiating array axioms, array operations and terms may be abstracted. The operations `select`, `store`, and `const-array` are mapped into corresponding uninterpreted functions, `select`, `store`, and `const-array` by extending the

EUF abstraction mapping  $\llbracket \cdot \rrbracket$  to array terms and operations as follows:

$$\llbracket a : \text{Array} \rrbracket = \hat{a} \quad (5)$$

$$\llbracket \text{select}(a, i) \rrbracket = \text{select}(\llbracket a \rrbracket, \llbracket i \rrbracket) \quad (6)$$

$$\llbracket \text{store}(a, i, x) \rrbracket = \text{store}(\llbracket a \rrbracket, \llbracket i \rrbracket, \llbracket x \rrbracket) \quad (7)$$

$$\llbracket \text{const-array}(k) \rrbracket = \text{const-array}(\llbracket k \rrbracket) \quad (8)$$

The array abstraction fits neatly into EUFORIA’s data abstraction approach. In fact, this abstraction approach keeps EUFORIA reasoning at the pure (quantifier-free) uninterpreted function level, for which there are efficient decision procedures. Consequently, EUFORIA’s EUF reachability algorithm works as-is; we discuss reachability in the next section.

#### A. Abstract IC3, or, The EUFORIA algorithm

EUFORIA is based on IC3 [2], a model checking algorithm for finite, Boolean transition systems. EUFORIA uses a counterexample-guided abstraction refinement (CEGAR) [8] approach that extends IC3 to apply to EUF transition systems while retaining termination.

EUFORIA takes a model checking problem as input,  $(X, Y, I, T, P)$ . It maps the CTS and property to produce a corresponding EUF abstract transition system (ATS) and property,  $(\hat{X}, \hat{Y}, \hat{I}, \hat{T}, \hat{P})$ . EUFORIA then alternates between two phases: EUF reachability and abstraction refinement. EUF reachability searches for a counterexample in the ATS. If no counterexample is found, soundness of the ATS proves that the property holds in the CTS. Otherwise, abstraction refinement analyzes the counterexample to determine if it is feasible in the CTS and, if not, modifies the EUF abstraction to increase its fidelity to the CTS. We first give a brief review of EUF reachability [1] before focusing on refinement.

As in IC3, the object for EUF reachability is an iteratively deepened sequence of reachable sets of formulas,  $R_i$ , each denoting an over-approximation of the set of states reachable in  $i$  transitions. The algorithm maintains the following invariants:

$$R_0 = \hat{I}(\hat{X}) \quad (9)$$

$$R_i \models R_{i+1} \quad (10)$$

$$R_i \models \hat{P}(\hat{X}) \quad (i < N) \quad (11)$$

$$R_{i+1} \text{ over-approximates the image of } R_i \quad (12)$$

EUF reachability computes a safety invariant for  $\hat{P}$  or a counterexample to the safety property. A safety invariant  $\hat{S}$  for  $\hat{P}$  has the following properties:

$$\hat{I} \models \hat{S}, \quad \hat{S} \wedge \hat{T} \models \hat{S}', \text{ and } \hat{S} \models \hat{P}.$$

#### B. Refinement

EUF reachability may find an abstract counterexample (ACX). Due to EUF abstraction, the concretized abstract counterexample (CACX) may *not* be a counterexample in the CTS. Consider the transition system  $\mathcal{E} = (X, Y, I, T)$  defined as:

$$(\{A, i\}, \emptyset, [\text{select}(A, i) = 3], [A' = \text{store}(A, i, 3)])$$

Consider the property,  $\text{select}(A, i) = 3$ , which is its own safety invariant. Nevertheless, it does not hold in  $\hat{\mathcal{E}}$ , since EUF abstraction does not preserve the relationship between store and select, and yields the CACX  $I \xrightarrow{T} \text{select}(A, i) \neq 3$  which is infeasible in the  $\text{QF\_ABV}$  theory. EUFORIA uses this contradictory CACX to *refine*, or increase the fidelity of, the abstraction. Refinement is accomplished by conjoining formulas, called lemmas, with the abstract transition relation.

In this example, EUFORIA learns an instance of McCarthy’s axiom (1), to eliminate the “spurious” violation caused by the abstraction:

$$\hat{A}' = \text{store}(\hat{A}, \hat{i}, \hat{3}) \Rightarrow \text{select}(\hat{A}', \hat{i}) = \hat{3}$$

This lemma constrains the abstract state space and is therefore appropriately called a *constraint lemma*. Constraint lemmas restrict the behavior of uninterpreted functions to make them conform more closely to the behavior of their concrete counterparts.

A second type of refinement involves learning new terms from CACXs. This kind of learning is similar to learning new predicates in a predicate abstraction (e.g., [9]). For instance,  $x' = x + 1 \wedge x' >_u x$  is invalid in  $\text{QF\_ABV}$  for the unsigned greater-than operator  $>_u$ , so EUFORIA can learn a lemma

$$\hat{x} \neq \widehat{2^{32} - 1} \vee \hat{x}' \neq \text{ADD}(\hat{x}, \hat{1}) \vee \neg \text{GT}(\hat{x}', \hat{x}).$$

This *expansion lemma* introduces a new term  $\widehat{2^{32} - 1}$  that increases the size of the abstract state space (the number of Herbrand models), making it more granular.

To sum up, constraint lemmas specialize UFs to particular concrete behaviors. Expansion lemmas increase the granularity of the EUF abstraction. EUFORIA learns array lemmas only if they crop up in a CACX’s contradiction, ensuring that the lemmas are directly relevant to the property that is being checked. Empirically speaking, contradictions usually feature a small handful of UFs which are ultimately relevant to the property, resulting in targeted lemmas. This process avoids most of the expense of array lemma generation, as we will see in the evaluation. We now describe how counterexample-guided learning is implemented in EUFORIA.

#### C. Implementation of Abstraction Refinement

An  $n$ -step *abstract counterexample* (ACX) is an execution  $\hat{A}_0 \xrightarrow{\hat{T}} \hat{A}_1 \xrightarrow{\hat{T}} \dots \xrightarrow{\hat{T}} \hat{A}_{n-1} \xrightarrow{\hat{T}} \hat{A}_n$  where each  $\hat{A}_i$  ( $0 \leq i \leq n$ ) is a state cube. An abstract formula  $\hat{\sigma}$  is *feasible* if its concretization  $\sigma$  is satisfiable over  $\text{QF\_ABV}$ ; therefore, an abstract counterexample is feasible if its concretization is a counterexample in the CTS.

EUFORIA’s refinement procedure, BUILD CX, is given in Figure 1; it has three stages. The first stage (lines 1–3) checks whether each  $\hat{A}_i$  is feasible ( $0 \leq i \leq n$ ). The second stage (lines 4–6) checks whether each  $\hat{A}_{i-1} \xrightarrow{\hat{T}} \hat{A}_i$  is feasible ( $0 < i \leq n$ ). The third stage, BUILD BMCCX, performs a bounded model checking (BMC) query to learn across multiple steps of the counterexample.

**BUILD\_CX():**

Returns true if counterexample is true, false if abstraction is refined

input:  $\hat{A}_0 \xrightarrow{\hat{T}} \hat{A}_1 \xrightarrow{\hat{T}} \dots \xrightarrow{\hat{T}} \hat{A}_{n-1} \xrightarrow{\hat{T}} \hat{A}_n$

- 1: **if**  $\exists i \in \{0, \dots, n\}. \neg \text{SAT}[A_i]$  **then**
- 2:    $\text{LEARN\_LEMMA}(\text{UNSAT\_CORE}())$
- 3:   **return** false
- 4: **if**  $\exists i \in \{1, \dots, n\}. \neg \text{SAT}[A_{i-1} \wedge T \wedge Y_{i-1} \wedge A_i]$  **then**
- 5:    $\text{LEARN\_LEMMA}(\text{UNSAT\_CORE}())$
- 6:   **return** false
- 7: **return** **BUILD\_BMC\_CX()**

Fig. 1. EUFORIA's refinement procedure

**BUILD\_BMC\_CX():**

- 1:  $\mathcal{B} \leftarrow \text{BMC\_FORMULA}()$
- 2: **if**  $\neg \text{SAT}[\mathcal{B}]$  **then**
- 3:    $\text{REFINE\_WITH\_INTERPOLANTS}(\text{UNSAT\_CORE}())$
- 4:   **return** false
- 5: **return** true  $\triangleright$  feasible counterexample

Fig. 2. Performs a bounded model check of the concretized counterexample

**REFINE\_WITH\_INTERPOLANTS(*core*):**

- 1:  $\mathcal{B}_{HC} \leftarrow \text{BUILD\_HORN}(\text{core})$
- 2:  $\mathcal{M} \leftarrow \text{HORN\_SOLVE}(\mathcal{B}_{HC})$
- 3: **for**  $i \in \{1, \dots, n\}$  **do**
- 4:    $p_i \leftarrow \text{GET\_INTERPOLANT}(\mathcal{M}, i)$
- 5:    $p_{i+1} \leftarrow \text{GET\_INTERPOLANT}(\mathcal{M}, i + 1)$
- 6:    $l \leftarrow p_{i-1}(X) \wedge \text{body}_i(X, Y, X') \wedge \neg p_i(X')$
- 7:    $\text{LEARN\_LEMMA}(l)$

Fig. 3. Constructs lemmas from an inductive interpolant sequence derived from a solution to (satisfiable) Horn clauses.  $\text{GET\_INTERPOLANT}(\mathcal{M}, i)$  returns a formula, the  $i$ th interpolant in the interpolant sequence, given a model for  $\mathcal{B}_{HC}$ .

**LEARN\_LEMMA(*f*):**

Precondition:  $f$  is unsatisfiable in  $\text{QF\_ABV}$

- 1:  $\hat{f} \leftarrow \text{ABSTRACT\_AND\_NORMALIZE}(f)$   $\triangleright$  abstract and eliminate input variables
- 2: **if**  $\hat{f}$  contains no inputs **then**
- 3:   **if**  $\text{VARS}(\hat{f}) \subseteq X$  **then**  $\triangleright$  only present-state vars
- 4:     Simplify and add lemma  $\neg \hat{f}(\hat{X}')$
- 5:   **if**  $\text{VARS}(\hat{f}) \subseteq X'$  **then**  $\triangleright$  only next-state vars
- 6:     Simplify and add lemma  $\neg \hat{f}(\hat{X})$
- 7: Simplify and add lemma  $\neg \hat{f}$

Fig. 4. Learns a lemma by abstracting and conjoining  $\hat{f}$  to  $\hat{T}$

If an infeasible state or transition is found during the first two stages, we compute an UNSAT core, negate it, and abstract it to form a constraint lemma (in  $\text{LEARN\_LEMMA}$ ). States and transitions are prioritized over BMC because it is advantageous to learn constraint lemmas, since they make the abstract state space smaller. Nevertheless, expansion lemmas are required to converge in general.

EUFORIA's expansion lemmas are computed in  $\text{BUILD\_BMC\_CX}$  with the help of interpolants that explain infeasible counterexamples. This stage of refinement has two phases. In phase one (lines 1–2),  $\text{BMC\_FORMULA}$  constructs the instance as below by explicitly renaming variables and using multiple copies of  $T$ :

$$\begin{aligned} \mathcal{B} = & A(X_0) \wedge T(X_0, Y_1, X_1) \wedge \\ & A(X_1) \wedge T(X_1, Y_2, X_2) \wedge \dots \wedge \\ & A(X_{n-1}) \wedge T(X_{n-1}, Y_n, X_n) \wedge A(X_n) \end{aligned}$$

If  $\mathcal{B}$  is feasible, it is a counterexample to the property.

If  $\mathcal{B}$  is infeasible,  $\text{BUILD\_BMC\_CX}$  enters phase two, implemented in  $\text{REFINE\_WITH\_INTERPOLANTS}$  (Figure 3).  $\text{BUILD\_HORN}$  uses  $\mathcal{B}$ 's UNSAT core to create a (reduced) inductive interpolant sequence problem  $\mathcal{B}_{HC}$  [10] using only the constraints from  $\mathcal{B}$  that occur in the core.  $\mathcal{B}_{HC}$  is a set of recursion-free Horn clauses in which uninterpreted predicates  $p_i$  stand for step-wise interpolants:

$$\begin{aligned} p_0(X_0) & \Leftarrow \text{true} \\ p_1(X_1) & \Leftarrow p_0(X_0) \wedge A^*(X_0) \wedge I(X_0) \wedge T^*(X_0, Y_1, X_1) \\ p_2(X_2) & \Leftarrow p_1(X_1) \wedge A^*(X_1) \wedge T^*(X_1, Y_2, X_2) \\ & \vdots \\ p_n(X_n) & \Leftarrow p_{n-1}(X_{n-1}) \wedge A^*(X_{n-1}) \wedge T^*(X_{n-1}, Y_n, X_n) \\ \text{false} & \Leftarrow p_n(X_n) \end{aligned}$$

where  $F^* = \bigwedge \{f \in F \mid f \in \text{UnsatCore}(\mathcal{B})\}$  for  $F \in \{A, T\}$ .<sup>1</sup> These Horn clauses are satisfiable by construction since  $\mathcal{B}$  is UNSAT.

For each nontrivial solution to the Horn clauses, we extract a lemma from the corresponding Horn clause as follows:

$$\neg[p_{i-1}(X) \wedge \text{body}_i(X, Y, X') \wedge \neg p_i(X')] \quad 0 < i \leq n \quad (13)$$

where  $\text{body}_i$  stands for the interpreted body predicates from the rule whose head is  $p_i$ . Lemmas are expansion lemmas only when the interpolants contain new terms. Using this method implies that the interpolation system itself decides whether a particular lemma is expansive or constraining (or both); EUFORIA does not make this decision explicitly. EUFORIA's back-end uses SPACER to solve  $\mathcal{B}_{HC}$ .

It is challenging to solve the BMC and interpolation queries for several reasons. First, there are multiple copies of  $T$ . Second,  $T$  is monolithic because it encodes the entire program, even though only a particular part of the program is relevant

<sup>1</sup> $\mathcal{B}_{HC}$  could be computed without  $\mathcal{B}$ 's UNSAT core, but using it promotes learning concise lemmas, because it substantially reduces the complexity of the Horn clause bodies. See equation (13).

for a given counterexample step. Third, even if we could reduce  $T$  at each step by removing irrelevant parts, using a large-step encoding [11] for  $T$  means that the reduced  $T$  would still likely contain a whole pile of nested Boolean logic, not all of which is necessarily relevant.

We address these difficulties by exploiting extra information available during abstract reachability. Whenever EUFORIA’s pre-image computation [1] generates a counterexample state  $\hat{A}_i$ , it also saves the abstract model of the abstract transition. This model is used to perform a model-based projection (called a model-based cone of influence traversal in [1]) of  $\hat{T}$  which produces extra constraints on  $T$  at each counterexample step. The BMC query is then pre-processed by an equation solving pass<sup>2</sup> that hoists equations from within the formula and substitutes them into the formula, removing many thousands of variables over multiple refinement passes.

This procedure prunes the search space significantly, addressing two of the three difficulties: the second ( $T$  is monolithic) and third ( $T$  transitions contain complex Boolean logic). The equation-solving pass is similar to STP’s [12]. We note, however, that equation-solving on its own is not enough; it needs the extra constraints on  $T$  to achieve peak efficiency.

#### D. Exceptionally Lazy Learning of Array Lemmas

The procedure `LEARNLEMMA` learns array lemmas. Fundamentally, it accomplishes this role by negating formulas found to be unsatisfiable in `QF_ABV`. The negated formulas are also simplified, specifically by possibly eliminating input variables, in order to generalize the lemmas as much as possible. Moreover, if the lemma formula is a state formula, then two versions are learned: one on current-state variables and one on next-state variables (lines 2–6, Figure 4).

Consequently, EUFORIA generates property-directed instantiations of array theory axioms. For instance, here is a lemma learned in one of our benchmarks:

$$A \neq \text{const-array}(0) \vee 0 \neq \text{select}(A, i) \quad (14)$$

This lemma is an instance of axiom (4). We also find instances of McCarthy’s axiom (1):

$$\text{select}(A', i) = 0 \vee i' \neq i \vee A' \neq \text{store}(A, i', 0) \quad (15)$$

Array lemmas may also include bit-vector function symbols to learn targeted lemmas about composite behavior:

$$B \neq \text{store}(A, i, 0) \vee \text{extract}(7, 0, \text{select}(B, i)) \neq 0 \quad (16)$$

Finally, some lemmas combine multiple array axioms:

$$\text{store}(B, i, 0) \neq A \vee \text{store}(A, i, 0) = A \quad (17)$$

This lemma relates stores and array extensionality. It is not a direct instance of any axiom (1)–(4), but rather a consequence of several instantiations.

<sup>2</sup>The `solve-eqs` tactic in Z3.

## IV. EVALUATION

To evaluate EUFORIA, we rely on benchmarks from SV-COMP’17 [13], as they are widely used and relatively well understood. Moreover, the benchmarks are categorized according to the types of properties, which gives us a way to validate our claims about the appropriateness of EUFORIA for control properties. We evaluate on C programs from the `Systems_DeviceDriversLinux64_ReachSafety` benchmark set, hereafter abbreviated *DeviceDrivers*. This set contains 64-bit C programs and contains “problems that require the analysis of pointer aliases and function pointers.”

We consider two other model checkers, SPACER and IC3IA. SPACER [14], [15], [16] is an over- and under-approximation driven incremental model checker that is tightly integrated with Z3. It computes procedure summaries to support checking programs with recursive functions. It is capable of inferring quantified array invariants and uses model-based projection array procedures to lazily instantiate property-directed array axioms, making the checker particularly efficient. IC3IA [9] is an IC3-style CEGAR model checker that implements implicit predicate abstraction. IC3IA’s architecture is quite similar EUFORIA’s, more similar than SPACER’s. As discussed in Cimatti [9], IC3IA is superior to state-of-the-art bit-level IC3 implementations and can support hundreds of predicates, around an order of magnitude more than what explicit predicate abstraction tools practically support. We also evaluated ELDARICA [17], a predicate-abstraction based CEGAR model checker that supports integers, algebraic data types, arrays, and bit vectors. Unfortunately, ELDARICA either threw errors, ran out of time, or ran out of memory on all of our benchmarks, so we do not consider it further.

We use SeaHorn as a front-end to encode programs into Horn clauses. SeaHorn [18] is a verification condition (VC) generator for C and C++ programs that uses LLVM in order to optimize and generate large-step, Horn clause benchmarks in SMT-LIBv2 `declare-rel` format. Note that we use the term benchmark to refer both to the C programs and their encoded counterparts. Since SeaHorn is not able to produce bit-vector encoded benchmarks, we modified it to produce bit-vector VCs. Moreover, since EUFORIA does not yet support procedure calls, we instruct SeaHorn to inline all procedures, resulting in linear Horn clauses. We ran SeaHorn on each benchmark, limiting it to one hour of runtime and 8GB of memory. SeaHorn can fail to produce a usable benchmark due to lack of resources or because the input is trivially solved during optimization. All told, SeaHorn produced 948 *DeviceDrivers* Horn clause benchmarks out of 2703 original C programs. 687 are safe and 261 are unsafe.

SPACER natively supports Horn clauses, but EUFORIA and IC3IA take VMT files as input. The VMT format [19] is a syntax-compatible extension of the SMT-LIBv2 format that fixes a syntax for labeling formulas denoting initial state, the transition relation, and property. In order to create comparable benchmarks for EUFORIA and IC3IA, we translate the Horn clause benchmarks into VMT using Horn2VMT [20], resulting




Fig. 5. Note the points on the right hand side of this plot. Each point is a benchmark that EUFORIA solved within 30 minutes that SPACER did not solve during that time.

in 948 VMT files that correspond to the 948 Horn benchmarks. The benchmarks range in size from  $2^9$  to more than  $2^{23}$ , with a median size of  $2^{19}$ ; this size is the number of distinct SMT-LIB expressions used to define  $(I, T, P)$ . When compressed with gzip, their sizes range from 2K to 153 MB.

All checkers run on 2.6 GHz Intel Sandy Bridge (Xeon E5-2670) machines with 2 sockets, 8 cores with 64GB RAM, running RedHat Enterprise Linux 7. Each checker run was assigned to one socket during execution and was given a 30 minute timeout. For every benchmark solved by any checker, we verified that its result was consistent with other checkers.

#### A. EUFORIA compared with SPACER

Figure 5 shows a scatter plot of runtime for EUFORIA and SPACER on *DeviceDrivers* benchmarks. Overall, EUFORIA solves 491 benchmarks and SPACER solves 386. EUFORIA times out on 33 benchmarks that SPACER solves. SPACER times out on 138 benchmarks that EUFORIA solves.

a) *When SPACER solves EUFORIA's timeouts:* In the 34 cases where spacer was able to solve a benchmark that EUFORIA could not, we identified several causes:

- 1) A number of benchmarks are quite large, and the overhead of a monolithic transition relation can dominate EUFORIA's abstract reachability. To explain: SeaHorn produces an *explicitly sliced* transition relation which SPACER exploits by making sliced incremental queries. EUFORIA consumes and queries a monolithic transition relation  $T$ , which is what Horn2VMT produces.
- 2) EUFORIA's front-end takes excessive time to parse and normalize some benchmarks. EUFORIA parses VMT

files using MathSAT5, since it the simplest API to do so. In addition to parsing, MathSAT normalizes and simplifies the resulting formula. For some benchmarks, these steps take longer than solving the benchmark itself. SPACER's parsing & simplification steps, on the other hand, are much quicker.

- 3) SPACER's preprocessor is able to solve some benchmarks without even invoking search.
- 4) EUFORIA gets stuck while computing interpolants. This is unexpected as EUFORIA uses SPACER to compute interpolants.

We believe that front-end improvements would address the issues identified in items 1–3. For instance, SPACER's preprocessor could be made independent of Z3 so that it could be applied before Horn2VMT.<sup>3</sup> Alternatively, EUFORIA could be integrated into Z3 so that it could exploit the same preprocessing as SPACER, but exploring this remains future work.

b) *When EUFORIA solves SPACER's timeouts:* In the 137 cases where EUFORIA was able to solve a benchmark that SPACER did not, we examined causes. In over half of the cases, SPACER gets stuck solving concrete incremental queries. In the other 52 cases, SPACER gives up before the timeout (it returns unknown). In other words, in every case individual queries were unable to be tackled given the resources constraints. Therefore we emphasize that, in contrast, EUFORIA has the strong benefit of making individual queries predictably fast.

We wondered: is EUFORIA only winning because it hardly needs to do refinement? The answer is no. Figure 6 shows the same scatter plot as Figure 5 but restricted to EUFORIA-solved benchmarks that required at least one abstraction refinement. It shows that EUFORIA requires refinement for many of the benchmarks for which SPACER times out.

#### B. EUFORIA compared with IC3IA

Figure 7 shows a scatter plot of our results compared with IC3IA. IC3IA solves 128 benchmarks total. Excepting 2 of these, EUFORIA solves *all* the benchmarks that IC3IA solves, usually in orders of magnitude less time. Our results are significant because IC3IA and EUFORIA are quite similar: both implement a PDR-style [21] algorithm, both operate on *exactly* the same VMT instance encoding, and both are written in C++. They differ in two respects: (1) IC3IA uses (implicit) predicate abstraction and EUFORIA uses EUF abstraction; (2) IC3IA's SMT solver backend is MathSAT5 and EUFORIA's is Z3.

On both of the benchmarks where EUFORIA doesn't terminate, gets stuck after several seconds in an interpolant query (the second or fifth refinement).

#### C. EUFORIA and array abstraction

For solvers that use lazy theory lemma learning or a trigger-based saturation method [22], array lemmas will be learned in response to property-directed queries. Does EUFORIA's array abstraction really provide a benefit over such an approach?

<sup>3</sup>We tried dumping the benchmark after SPACER's preprocessing step, but the benchmark was no longer guaranteed to be Horn, so it was not a valid input for encoding to VMT with Horn2VMT.




Fig. 6. EUFORIA vs SPACER restricted to those benchmarks EUFORIA solves that require *at least one* abstraction refinement.

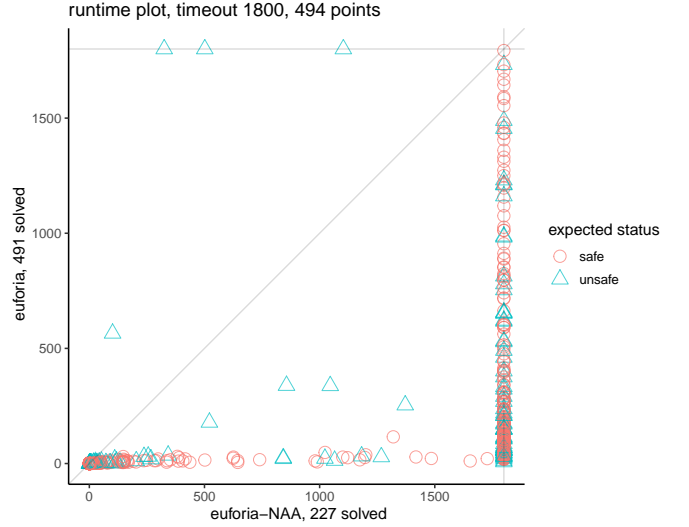


Fig. 8. EUFORIA<sub>NAA</sub> (no array abstraction) (x axis) compared with EUFORIA (y axis).




Fig. 7. EUFORIA vs IC3IA

To address this question, we modified EUFORIA to compute a hybrid abstraction using the theory of EUF and arrays. It abstracts bit-vector operations into UFs (as before), but uses array theory operations for arrays. Call this configuration EUFORIA<sub>NAA</sub>, for No Array Abstraction.

As demonstrated in Figure 8, EUFORIA<sub>NAA</sub> is significantly slower almost everywhere and strictly slower in all cases but one. One important difference between EUFORIA and EUFORIA<sub>NAA</sub> is an enormous disparity in array theory lemmas learned by the underlying SMT solver. Between configurations, the difference of the number of array theory lemma instantiations is almost two orders of magnitude (1.9), on 95% of the benchmarks; almost four orders of magnitude (3.8), on 50% of the benchmarks; and more than seven orders of magnitude (7.2), on 5%. To calculate this result, we measure the number of array theory axiom instantiations in the underlying SMT solver (Z3). Then, for each benchmark, we took the difference of the logs (base 10) between the two configurations; this quantity is proportional to the order of magnitude difference between the numbers.

We conclude that EUFORIA<sub>NAA</sub> spends a lot of time reasoning about arrays despite the fact that EUFORIA required relatively little array reasoning to solve the same benchmarks. Moreover, compared to SPACER’s 386 solves, EUFORIA<sub>NAA</sub> solves only 227 instances, which (1) shows that array abstraction is critical to performance and (2) gives some additional evidence that SPACER’s array projection helps its runtime.

#### D. EUFORIA in itself—the role of lemmas

This section discusses EUFORIA’s learned lemmas as detailed in Section III-B. Lemmas in general play a relatively






Fig. 9. Breakdown of lemmas as array-related and non array-related on the subset of benchmarks (92) for which any lemma learning was required.

minor role; they’re only required in 19% of euforia’s solved benchmarks (92). Figure 9 shows the count of total lemmas learned for each such benchmark, broken down by whether EUFORIA learned array lemmas or non-array lemmas. First, we can see that there is a trend that EUFORIA learns fewer array lemmas than data lemmas. Second, all but two benchmarks required fewer than 100 lemmas. These results suggest our benchmarks only depend sparingly on the behavior of memory manipulations, giving some validation that EUFORIA’s abstraction works well for control-dominated benchmarks.

## V. RELATED WORK

The relationship between EUF and the theory of arrays has been long recognized [23], [24] and analyzed [25] and exploited in decision procedures [26] and in the implementation of several SMT solvers, including Yices [27] and Z3 [28]. Array terms are compiled into EUF or a ground theory to instantiate the needed array axioms. Our approach lifts EUF outside the SMT solver, to the model checking level, and refines it on demand.

Komuravelli *et al.* introduce a model-based projection for pre-images in order to rewrite array operators into terms in a scalar theory [16]; this algorithm is implemented in SPACER [15] used in our evaluation. Predicate abstraction applies to programs with arrays directly [9], with the limitation that quantifier-free interpolants do not exist in general for the theory of arrays [24]. We inherit that limitation, but contribute a different, inexpensive way to place array constraints in pre-images and refine them lazily.

Broadly, SMT solvers solve constraints over arrays in three ways (sometimes combined): (1) by rewriting selects and

stores into a finite number of terms and axiom instantiations in a ground theory, possibly combined with EUF [29], [30], [23], [31], [25], [32], [33], [22], [26]; (2) by abstraction-refinement procedures over the array constraints [12], [34]; (3) by rewriting into (non-abstract) representations which are solved with specialized algorithms [35], [36], [37]. The issue addressed by our paper is applicable to each of these: we use an abstraction that inexpensively supports (limited) array reasoning and we only invoke an SMT array solver at the last possible moment.

## VI. CONCLUSION AND FUTURE WORK

This paper introduces an approach for model checking software with arrays that avoids substantial computational effort spent in reasoning about arrays by using EUF abstraction. We integrated our approach inside an incremental model checker that natively supports EUF abstraction. Our approach bests stiff competition on control-oriented benchmarks, solving over 100 more benchmarks.

We demonstrated that our approach reduces the amount of redundant or irrelevant array reasons by several orders of magnitude in most cases. We are eager to investigate the possibilities of expanding our universe of target programs. As software size grows, its sheer size begins to overwhelm the checker, even if the property to prove is relatively simple (for a machine). Inlining all functions only exacerbates the problem. In future work we plan to explore compositional reasoning, in particular analyzing programs with procedures by integrating it efficiently with our EUF abstraction.

We find that for some benchmarks, even for control properties, stronger lemmas are required to aid convergence. We would like to address this by inferring quantified lemmas during search. One issue is how to generalize counterexamples to quantified lemmas. A second issue is how to keep the abstraction tractable in the presence of quantified lemmas. Both of these issues form important future work.

## ACKNOWLEDGMENT

We thank Arie Gurfinkel and Nikolaj Bjørner for their help regarding SPACER and Z3 internals.

## REFERENCES

- [1] D. Bueno and K. A. Sakallah, “EUFORIA: Complete software model checking with uninterpreted functions,” in *Verification, Model Checking, and Abstract Interpretation - 20th International Conference, VMCAI 2019, Cascais, Portugal, January 13-15, 2019, Proceedings*, ser. Lecture Notes in Computer Science, C. Enea and R. Piskac, Eds., vol. 11388. Springer, 2019, pp. 363–385.
- [2] A. R. Bradley, “SAT-based model checking without unrolling,” in *Verification, Model Checking, and Abstract Interpretation*, ser. Lecture Notes in Computer Science, R. Jhala and D. A. Schmidt, Eds., vol. 6538. Springer, 2011, pp. 70–87.
- [3] J. McCarthy, “Towards a mathematical science of computation,” in *Information Processing, Proceedings of the 2nd IFIP Congress 1962, Munich, Germany, August 27 - September 1, 1962*. North-Holland, 1962, pp. 21–28.
- [4] E. M. Clarke, O. Grumberg, and D. E. Long, “Model checking and abstraction,” *ACM Trans. Program. Lang. Syst.*, vol. 16, no. 5, pp. 1512–1542, 1994.



- [5] A. R. Bradley and Z. Manna, "Checking safety by inductive generalization of counterexamples to induction," in *Formal Methods in Computer-Aided Design*. IEEE Computer Society, 2007, pp. 173–180.
- [6] C. Barrett, A. Stump, and C. Tinelli, "The SMT-LIB Standard: Version 2.0," in *Workshop on Satisfiability Modulo Theories*, A. Gupta and D. Kroening, Eds., 2010.
- [7] J. R. Burch and D. L. Dill, "Automatic verification of pipelined microprocessor control," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, D. L. Dill, Ed., vol. 818. Springer, 1994, pp. 68–80.
- [8] E. M. Clarke, O. Grumberg, S. Jha, Y. Lu, and H. Veith, "Counterexample-guided abstraction refinement," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, E. A. Emerson and A. P. Sistla, Eds., vol. 1855. Springer, 2000, pp. 154–169.
- [9] A. Cimatti, A. Griggio, S. Mover, and S. Tonetta, "IC3 modulo theories via implicit predicate abstraction," in *Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, E. Ábrahám and K. Havelund, Eds., vol. 8413. Springer, 2014, pp. 46–61.
- [10] P. Rümmer, H. Hojjat, and V. Kuncak, "Classifying and solving horn clauses for verification," in *Verified Software: Theories, Tools, Experiments*, ser. Lecture Notes in Computer Science, E. Cohen and A. Rybalchenko, Eds., vol. 8164. Springer, 2013, pp. 1–21. [Online]. Available: [https://doi.org/10.1007/978-3-642-54108-7\\_1](https://doi.org/10.1007/978-3-642-54108-7_1)
- [11] D. Beyer, A. Cimatti, A. Griggio, M. E. Keremoglu, and R. Sebastiani, "Software model checking via large-block encoding," in *Formal Methods in Computer-Aided Design*. IEEE, 2009, pp. 25–32.
- [12] V. Ganesh and D. L. Dill, "A decision procedure for bit-vectors and arrays," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, W. Damm and H. Hermanns, Eds., vol. 4590. Springer, 2007, pp. 519–531.
- [13] D. Beyer, "Software verification with validation of results - (report on SV-COMP 2017)," in *Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, A. Legay and T. Margaria, Eds., vol. 10206, 2017, pp. 331–349.
- [14] A. Komuravelli, A. Gurfinkel, S. Chaki, and E. M. Clarke, "Automatic abstraction in smt-based unbounded software model checking," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, N. Sharygina and H. Veith, Eds., vol. 8044. Springer, 2013, pp. 846–862. [Online]. Available: [https://doi.org/10.1007/978-3-642-39799-8\\_59](https://doi.org/10.1007/978-3-642-39799-8_59)
- [15] A. Komuravelli, A. Gurfinkel, and S. Chaki, "SMT-Based Model Checking for Recursive Programs," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, A. Biere and R. Bloem, Eds., vol. 8559. Berlin, Heidelberg: Springer-Verlag, 2014, pp. 17–34. [Online]. Available: <https://doi.org/10.1007/978-3-319-08867-9>
- [16] A. Komuravelli, N. Bjørner, A. Gurfinkel, and K. L. McMillan, "Compositional verification of procedural programs using horn clauses over integers and arrays," in *Formal Methods in Computer-Aided Design*, R. Kaivola and T. Wahl, Eds. IEEE, 2015, pp. 89–96.
- [17] H. Hojjat and P. Rümmer, "The ELDARICA horn solver," in *2018 Formal Methods in Computer Aided Design, FMCAD 2018, Austin, TX, USA, October 30 - November 2, 2018*, N. Bjørner and A. Gurfinkel, Eds. IEEE, 2018, pp. 1–7. [Online]. Available: <https://doi.org/10.23919/FMCAD.2018.8603013>
- [18] A. Gurfinkel, T. Kahsai, A. Komuravelli, and J. A. Navas, "The SeaHorn verification framework," in *Computer Aided Verification*, ser. Lecture Notes in Computer Science, D. Kroening and C. S. Pasareanu, Eds., vol. 9206. Springer, 2015, pp. 343–361.
- [19] F. B. Kessler, *The VMT format*, 2020 (accessed May 13, 2020). [Online]. Available: <https://nuxmv.fbk.eu/index.php?n=Languages.VMT>
- [20] D. Bueno and K. Sakallah, "Horn2VMT: Translating horn reachability into transition systems," in *Workshop on Horn Clauses for Verification and Synthesis*, 2020, p. To appear.
- [21] N. Een, A. Mishchenko, and R. Brayton, "Efficient implementation of property directed reachability," in *Formal Methods in Computer-Aided Design*. IEEE, 2011, pp. 125–134.
- [22] L. M. de Moura and N. Bjørner, "Generalized, efficient array decision procedures," in *Formal Methods in Computer-Aided Design*. IEEE, 2009, pp. 45–52. [Online]. Available: <https://doi.org/10.1109/FMCAD.2009.5351142>
- [23] D. Kapur and C. G. Zarba, "A reduction approach to decision procedures," University of New Mexico, Tech. Rep., 2005.
- [24] D. Kapur, R. Majumdar, and C. G. Zarba, "Interpolation for data structures," in *SIGSOFT FSE*, M. Young and P. T. Devanbu, Eds. ACM, 2006, pp. 105–116. [Online]. Available: <https://doi.org/10.1145/1181775.1181789>
- [25] A. Goel, S. Krstić, and A. Fuchs, "Deciding array formulas with frugal axiom instantiation," in *International Workshop on Satisfiability Modulo Theories*, ser. SMT '08. New York, NY, USA: ACM, 2008, pp. 12–17. [Online]. Available: <http://doi.acm.org/10.1145/1512464.1512468>
- [26] A. R. Bradley, Z. Manna, and H. B. Sipma, "What's decidable about arrays?" in *Verification, Model Checking, and Abstract Interpretation*, ser. Lecture Notes in Computer Science, E. A. Emerson and K. S. Namjoshi, Eds., vol. 3855. Springer, 2006, pp. 427–442. [Online]. Available: [https://doi.org/10.1007/11609773\\_28](https://doi.org/10.1007/11609773_28)
- [27] B. Dutertre and L. M. de Moura, "A fast linear-arithmetic solver for DPLL(T)," in *Computer Aided Verification, 18th International Conference, CAV 2006, Seattle, WA, USA, August 17-20, 2006, Proceedings*, ser. Lecture Notes in Computer Science, T. Ball and R. B. Jones, Eds., vol. 4144. Springer, 2006, pp. 81–94. [Online]. Available: [https://doi.org/10.1007/11817963\\_11](https://doi.org/10.1007/11817963_11)
- [28] L. M. de Moura and N. Bjørner, "Z3: An efficient SMT solver," in *Tools and Algorithms for the Construction and Analysis of Systems*, ser. Lecture Notes in Computer Science, C. R. Ramakrishnan and J. Rehof, Eds., vol. 4963. Springer, 2008, pp. 337–340.
- [29] N. Suzuki and D. Jefferson, "Verification decidability of presburger array programs," CARNEGIE-MELLON UNIV PITTSBURGH PA DEPT OF COMPUTER SCIENCE, Tech. Rep., 1977.
- [30] J. Jaffar, "Presburger arithmetic with array segments," *Inf. Process. Lett.*, vol. 12, no. 2, pp. 79–82, 1981. [Online]. Available: [https://doi.org/10.1016/0020-0190\(81\)90007-7](https://doi.org/10.1016/0020-0190(81)90007-7)
- [31] C. Lynch and B. Morawska, "Automatic decidability," in *IEEE Symposium on Logic in Computer Science*. IEEE Computer Society, 2002, p. 7. [Online]. Available: <https://doi.org/10.1109/LICS.2002.1029813>
- [32] A. Armando, M. P. Bonacina, S. Ranise, and S. Schulz, "New results on rewrite-based satisfiability procedures," *ACM Trans. Comput. Log.*, vol. 10, no. 1, pp. 4:1–4:51, 2009. [Online]. Available: <https://doi.org/10.1145/1459010.1459014>
- [33] J. Christ and J. Hoenicke, "Weakly equivalent arrays," in *International Symposium on Frontiers of Combining Systems*, ser. FroCoS 2015. Berlin, Heidelberg: Springer-Verlag, 2015, pp. 119–134.
- [34] R. Brummayer and A. Biere, "Lemmas on demand for the extensional theory of arrays," *JSAT*, vol. 6, no. 1-3, pp. 165–201, 2009. [Online]. Available: <https://satassociation.org/jsat/index.php/jsat/article/view/74>
- [35] A. Stump, C. W. Barrett, D. L. Dill, and J. R. Levi, "A decision procedure for an extensional theory of arrays," in *IEEE Symposium on Logic in Computer Science*. IEEE Computer Society, 2001, pp. 29–37.
- [36] M. Bofill, R. Nieuwenhuis, A. Oliveras, E. Rodríguez-Carbonell, and A. Rubio, "A write-based solver for SAT modulo the theory of arrays," in *Formal Methods in Computer-Aided Design*, A. Cimatti and R. B. Jones, Eds. IEEE, 2008, pp. 1–8. [Online]. Available: <https://doi.org/10.1109/FMCAD.2008.ECP18>
- [37] P. Habermehl, R. Iosif, and T. Vojnar, "What else is decidable about integer arrays?" in *Foundations of Software Science and Computational Structures*, ser. Lecture Notes in Computer Science, R. M. Amadio, Ed., vol. 4962. Springer, 2008, pp. 474–489.
- [38] *Formal Methods in Computer-Aided Design*. IEEE, 2009.